

# CSE 120/220: Operating Systems Principles

## Lecture 5: Synchronization

Prof. J. Pasquale  
University of California, San Diego  
January 27, 2025

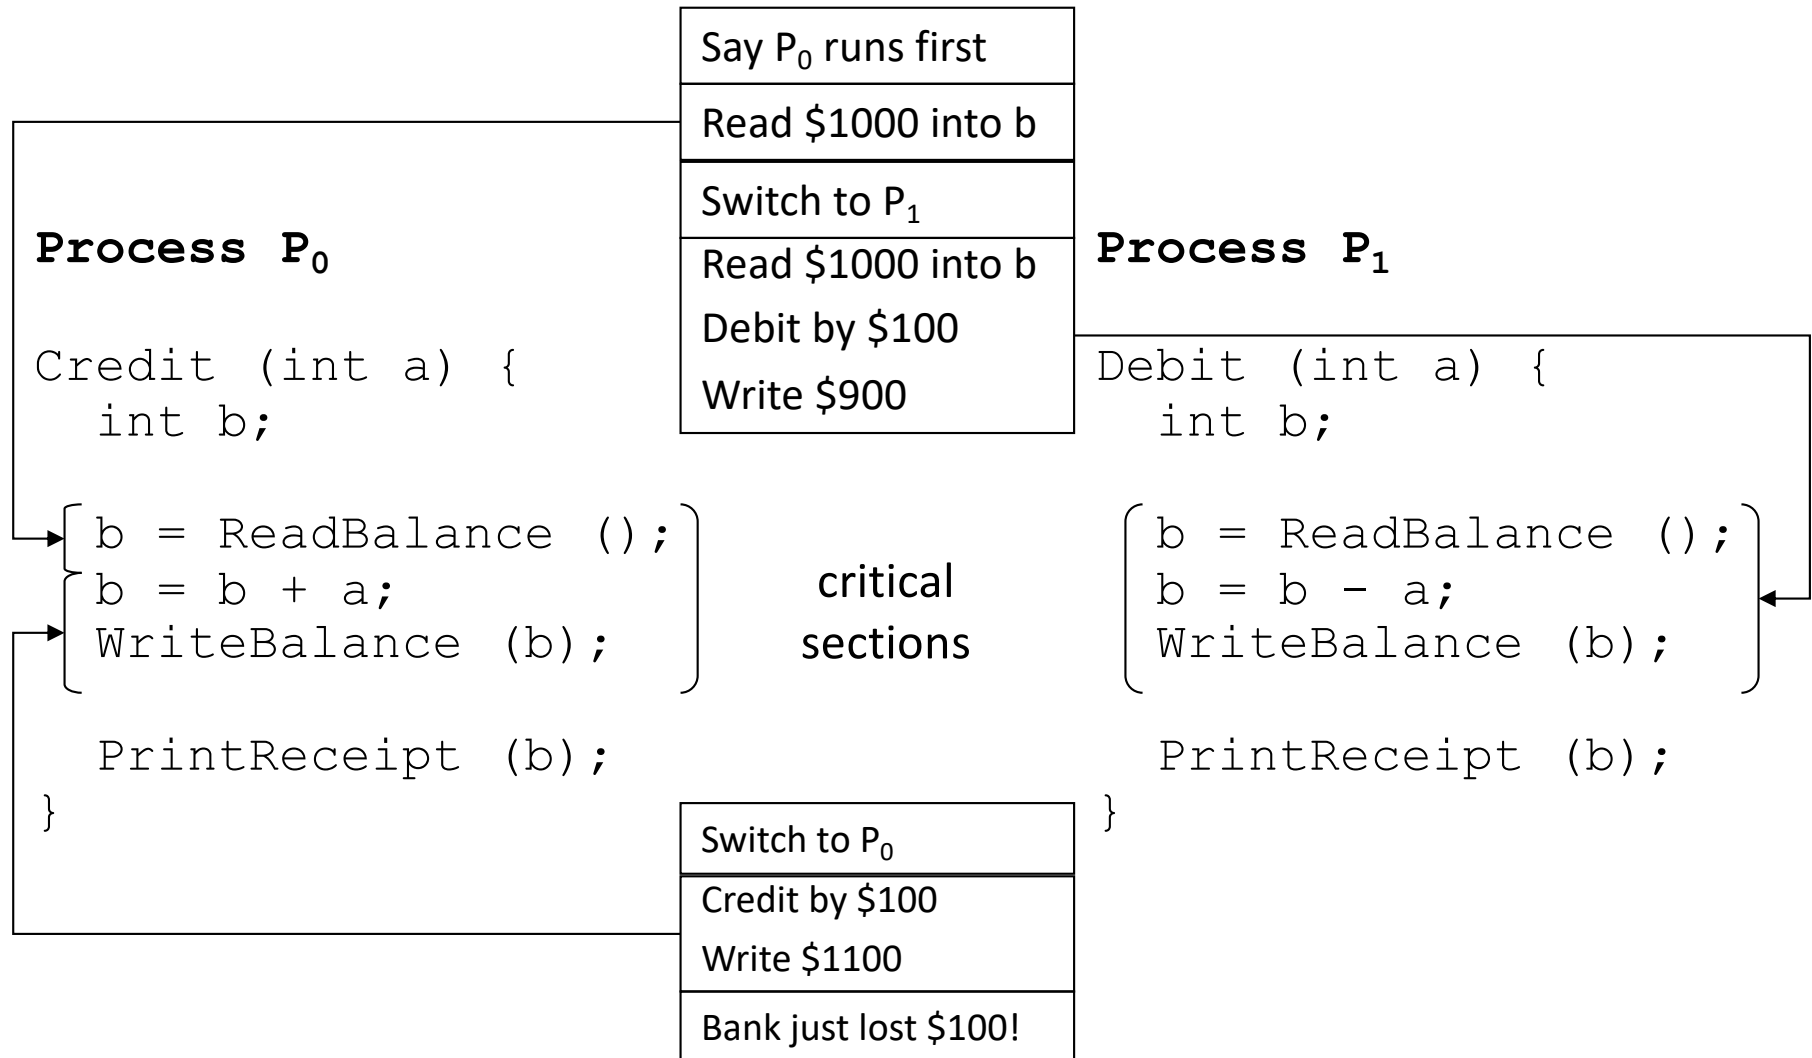
# Synchronization

- Synchronize: events happen at the same time
- Process synchronization
  - Events in processes that occur “at the same time”
  - Actually, when one process waits for another
- Uses of synchronization
  - Prevent race conditions
  - Wait for resources to become available


# The Credit/Debit Problem

- Say you have \$1000 in your bank account
- You deposit \$100
- You also withdraw \$100
- How much should be in your account?
- What if deposit/withdraw occur at same time?

# Credit/Debit Problem: Race Condition



# To Avoid Race Conditions




- Identify related *critical sections*
  - Section(s) of code executed by different processes
  - Must run *atomically, with respect to each other*
- Enforce mutual exclusion
  - Only one process active in a critical section

# What Does Atomic Really Mean?

- Atomic means “indivisible”
- We seek *effective* atomicity
  - can interrupt, as long as interruption has no effect
- It *is* OK to interrupt process in critical section
  - as long as other processes have no effect
- How to determine
  - Consider effect of critical section in isolation
  - Next consider interruptions: if same result, OK

# How to Achieve Mutual Exclusion?



- Surround critical section with entry/exit code
- Entry code should act as a barrier
  - If another process is in critical section, block
  - Otherwise, allow process to proceed
- Exit code should release other entry barriers

# Requirements for Good Solution

- Given multiple cooperating processes
    - Each process has a critical section
    - All critical sections are to be mutually exclusive
1. At most one in a critical section at a time
  2. Can't prevent entry if no others are in theirs
  3. Should eventually be able to enter
  4. No assumptions about CPU speed or number



# Software Lock?

```
shared int lock = OPEN;
```

**P<sub>0</sub>**

```
while (lock == CLOSED);
```

```
lock = CLOSED;
```

```
< critical section >
```

```
lock = OPEN;
```

**P<sub>1</sub>**

```
while (lock == CLOSED);
```

```
lock = CLOSED;
```

```
< critical section >
```

```
lock = OPEN;
```

- Lock indicates if any process in critical section
- Is there a problem?

# Take Turns?

```
shared int turn = 0;    // arbitrary set to  $P_0$ 
```

**$P_0$**

```
while (turn != 0);  
< critical section >  
turn = 1;
```

**$P_1$**

```
while (turn != 1);  
< critical section >  
turn = 0;
```

- Alternate which process enters critical section
- Is there a problem?

# State Intention?

```
shared boolean intent[2] = {FALSE, FALSE};
```

**P<sub>0</sub>**

```
intent[0] = TRUE;  
while (intent[1]);  
< critical section >  
intent[0] = FALSE;
```

**P<sub>1</sub>**

```
intent[1] = TRUE;  
while (intent[0]);  
< critical section >  
intent[1] = FALSE;
```

- Process states intent to enter critical section
- Is there a problem?

# Peterson's Solution

```
shared int turn;  
shared boolean intent[2] = {FALSE, FALSE};
```

**P<sub>0</sub>**

```
intent[0] = TRUE;  
turn = 1;  
while (intent[1] && turn==1);  
< critical section >
```

```
intent[0] = FALSE;
```

**P<sub>1</sub>**

```
intent[1] = TRUE;  
turn = 0;  
while (intent[0] && turn==0);  
< critical section >
```

```
intent[1] = FALSE;
```

- If competition, take turns; otherwise, enter
- Is there a problem?
- There is a version for  $n \geq 2$ ; more complex

# What about Disabling Interrupts?

- Reasoning
  - No interrupts  $\Rightarrow$  no uncontrolled context switches
  - No uncontrolled context switches  $\Rightarrow$  no races
  - No races  $\Rightarrow$  mutual exclusion
- Is there a problem?

# Test-and-Set Lock Instruction: TSL

- TSL mem (test-and-set lock: contents of mem)  
do atomically (i.e., locking the memory bus)  
[ **test** if mem == 0 AND **set** mem = 1 ]
- Operations occur without interruption
  - Memory bus is locked
  - Not affected by hardware interrupts

# What TSL Does, Expressed in C

- Assume C function, TSL(int \*), that is *atomic*

```
TSL(int *lockptr)
```

```
{  int oldval;
```

```
    oldval = *lockptr
```

```
    *lockptr = 1;
```

```
    return ((oldval == 0) ? 1 : 0);
```

```
}
```

Effectively, this is what the  
TSL instruction does in one  
atomic instruction

# Mutual Exclusion Using TSL

```
shared int lock = 0;
```

**P<sub>0</sub>**

```
while (! TSL(&lock));  
< critical section >  
lock = 0;
```

**P<sub>1</sub>**

```
while (! TSL(&lock));  
< critical section >  
lock = 0;
```

- Shared variable solution using TSL(int \*)
  - tests if lock == 0 (if so, will return 1; else 0)
  - before returning, sets lock to 1
- Simple, works for any number of processes
- Still “suffers” from busy waiting



# Semaphores

- Synchronization variable
  - Takes on integer values
  - Can cause a process to block/unblock
- wait and signal operations
  - wait (s)    decrement; block if  $< 0$
  - signal (s)    increment; if any blocked, unblock one
- No other operations allowed
  - In particular, cannot test value of semaphore!

# Examples and Interpretation

- wait (s)      decrement; block if  $< 0$
- signal (s)    increment; if any blocked, unblock
  
- wait (1)       $s \rightarrow 0$     GO
- wait (0)       $s \rightarrow -1$    STOP (i.e., block)
- signal (-1)    $s \rightarrow 0$     GO *and* allow one to GO
- signal (0)     $s \rightarrow 1$     GO

# Mutual Exclusion

```
sem mutex = 1;           // declare and initialize
```

**P<sub>0</sub>**

```
wait (mutex);  
< critical section >  
signal (mutex);
```

**P<sub>1</sub>**

```
wait (mutex);  
< critical section >  
signal (mutex);
```

- Use “mutex” semaphore, initialized to 1
- Only one process can enter critical section
- Simple, works for  $n$  processes
- Is there any busy-waiting?

# Order How Processes Execute

```
sem cond = 0;
```

**P<sub>0</sub>**

< to be done before P<sub>1</sub> >

```
signal (cond);
```

**P<sub>1</sub>**

```
wait (cond);
```

< to be done after P<sub>0</sub> >

- Cause a process to wait for another
- Use semaphore indicating condition; initially 0
  - the condition in this case: “P<sub>0</sub> has completed”
- Used for ordering processes
  - In contrast to mutual exclusion

# Semaphores: Only Synchronization

- Semaphores only provide synchronization
  - Synchronization: when a process blocks for event
- But, no information transfer
  - No way for a process to tell it blocked

# Semaphore Implementation

- Semaphore  $s = [n, L]$ 
  - $n$ : takes on integer values
  - $L$ : list of processes blocked on  $s$

- Operations

```
wait (sem s) {  
    s.n = s.n - 1;  
    if (s.n < 0) add calling process to S.L and block; }  
  
signal (sem s) {  
    s.n = s.n + 1;  
    if (s.L !empty) remove/unblock a process from s.L; }
```

# Alternative Implementation

- Semaphore  $s = [n, L]$ 
  - $n$ : takes on integer values, non-negative
  - $L$ : list of processes blocked on  $s$

- Operations

```
wait (sem s) {  
    if (s.n == 0) add calling process to s.L and block;  
    else s.n = s.n - 1; }  
  
signal (sem s) {  
    if (s.L !empty) remove/unblock a process from s.L;  
    else s.n = s.n + 1; }
```

# Wait and Signal Must Be Atomic

- Bodies of wait and signal are critical sections
- So, still need mechanism for mutual exclusion!
- Use a lower-level (more basic) mechanism
  - Test-and-set lock
  - Peterson's solution
- So, busy-waiting still exists (can never remove)
  - But at lower-level (within semaphore operations)
  - Occurrence limited to brief/known periods of time



# Analysis: Lower-Level Busy Waiting

- A calls wait (s), switch to B, B calls wait (s)
  - Switch occurs while A executing body of wait
- Body of wait is critical section, so B must block
  - Use test-set lock or Peterson's: busy waiting
- How long will B be blocked?
  - For time it takes to execute body of wait
- Small/known amount of time!
  - Compare to user critical section: unknown time

# Are These Equivalent?

## Implementation 1

```
wait (sem s) {  
    s.n = s.n - 1;  
    if s.n < 0 {  
        addProc (me, s.L);  
        block (me);  
    }  
}  
  
signal (sem s) {  
    s.n = s.n + 1;  
    if (! empty (s.L)) {  
        p = removeProc (s.L);  
        unblock (p);  
    }  
}
```

## Implementation 2

```
wait (sem s) {  
    if s.n ≤ 0 {  
        addProc (me, s.L);  
        block (me);  
    }  
    s.n = s.n - 1;  
}  
  
signal (sem s) { // same  
    s.n = s.n + 1;  
    if (! empty (s.L)) {  
        p = removeProc (s.L);  
        unblock (p);  
    }  
}
```

# Summary

- Synchronization: process waiting for another
- Critical section: code allowing race condition
- Mutual exclusion: one process excludes others
- Mutual exclusion mechanism: obey four rules
- Peterson's solution: all software, but complex
- Semaphores: simple flexible synchronization
  - wait and signal must be atomic, thus requiring lower-level mutual exclusion (Peterson's, TSL)

# Textbook

- OSP: Chapter 5
- OSC: Chapter 5 (Process Synchronization)
  - Lecture-related: 6.1-6.6

# Supplementary

- For those who wish to understand more subtle and advanced issues
- Will not be on exams

# Mutual Exclusion Using TSL

- Critical section entry code

```
                                ; assume lock initially 0
loop:  TSL REG, lock            ; atomically {load REG with lock
                                ;           and store 1 into lock}
                                ;
                                ; is REG (was lock) equal to 0?
                                ; if not equal to 0, check again
                                ; also known as a "spin lock"
                                ;
                                ;
```

- Critical section exit code

```
MOV lock, #0                    ; reset lock to 0
```

# Test and Test-and-Set

```
shared int lock = 0;
```

```
do {  
    while (lock == 1);           // cheap  
} while (! TSL(&lock));         // expensive  
< critical section >  
lock = 0;
```

- Busy-wait using simple reads of lock
  - Low overhead
- When lock opens, use test-and-set
  - Higher-overhead atomic operation less frequent

# Efficient No-Spin Locking Code

- Entry code

```
while (! TSL(&lock)) {           // lock closed
    yield ();                     // give up CPU
}
```

- Exit code

```
lock = 0;                        // open lock
```



# When is Busy-Waiting OK?

- Expected wait time < scheduling overhead
- Lots of processors (i.e., if waste is OK)
- Blocking is not an option (e.g., inside kernel)

# How Costly is Busy Waiting?

- Consider time spent in critical section
  - Chance of context switch increases with length
  - If switch to process seeking entry, it will busy wait
  - Wastes an entire quantum – this is cost
- So, try to minimize time in critical section
- Compare critical sections that are
  - user code (e.g., application code)
  - system code (e.g., semaphore operations)

# Making wait and signal Atomic

## Implementation 1

```
wait (sem s) {  
    s.n = s.n - 1;  
    if s.n < 0 {  
        addProc (me, s.L);  
        block (me);  
    }  
}
```

```
signal (sem s) {  
    s.n = s.n + 1;  
    if (! empty (s.L)) {  
        p = removeProc (s.L);  
        unblock (p);  
    }  
}
```

- To make atomic, add <entry> and <exit> code
- Where?

# Add Entry/Exit: Mutual Exclusion

## Implementation 1

```
wait (sem s) {  
    <entry>  
    s.n = s.n - 1;  
    if s.n < 0 {  
        addProc (me, s.L);  
        block (me);  
    }  
    <exit>  
}
```

```
signal (sem s) {  
    <entry>  
    s.n = s.n + 1;  
    if (! empty (s.L)) {  
        p = removeProc (s.L);  
        unblock (p);  
    }  
    <exit>  
}
```

# What Happens if block in wait?

## Implementation 1

```
wait (sem s) {  
    <entry>  
    s.n = s.n - 1;  
    if s.n < 0 {  
        addProc (me, s.L);  
        block (me);  
    }  
    <exit>  
}
```

```
signal (sem s) {  
    <entry>  
    s.n = s.n + 1;  
    if (! empty (s.L)) {  
        p = removeProc (s.L);  
        unblock (p);  
    }  
    <exit>  
}
```

- If wait blocks, how can signal ever run?
- Leads to deadlock

# No Mutual Exclusion During block

## Implementation 1

```
wait (sem s) {  
    <entry>  
    s.n = s.n - 1;  
    if s.n < 0 {  
        addProc (me, s.L);  
        <exit>  
        block (me);  
        <entry>  
    }  
    <exit>  
}  
  
signal (sem s) {  
    <entry>  
    s.n = s.n + 1;  
    if (! empty (s.L)) {  
        p = removeProc (s.L);  
        unblock (p);  
    }  
    <exit>  
}
```

- Give up mutual exclusion before blocking
- Reestablish after unblocking

# Making wait and signal Atomic

## Implementation 2

```
wait (sem s) {  
    if s.n ≤ 0 {  
        addProc (me, s.L);  
        block (me);  
    }  
    s.n = s.n - 1;  
}  
  
signal (sem s) {  
    s.n = s.n + 1;  
    if (! empty (s.L)) {  
        p = removeProc (s.L);  
        unblock (p);  
    }  
}
```

- To make atomic, add <entry> and <exit> code
- Where?

# Add Entry/Exit: Mutual Exclusion

## Implementation 2

```
wait (sem s) {  
    <entry>  
    if s.n ≤ 0 {  
        addProc (me, s.L);  
        block (me);  
    }  
    s.n = s.n - 1;  
    <exit>  
}
```

```
signal (sem s) {  
    <entry>  
    s.n = s.n + 1;  
    if (! empty (s.L)) {  
        p = removeProc (s.L);  
        unblock (p);  
    }  
    <exit>  
}
```



# What Happens if block in wait?

## Implementation 2

```
wait (sem s) {  
    <entry>  
    if s.n ≤ 0 {  
        addProc (me, s.L);  
        block (me);  
    }  
    s.n = s.n - 1;  
    <exit>  
}
```

```
signal (sem s) {  
    <entry>  
    s.n = s.n + 1;  
    if (! empty (s.L)) {  
        p = removeProc (s.L);  
        unblock (p);  
    }  
    <exit>  
}
```

- If wait blocks, how can signal ever run?

# No Mutual Exclusion During block

## Implementation 2

```
wait (sem s) {  
    <entry>  
    if s.n ≤ 0 {  
        addProc (me, s.L);  
        <exit>  
        block (me);  
        <entry>  
    }  
    s.n = s.n - 1;  
    <exit>  
}
```


```
signal (sem s) {  
    <entry>  
    s.n = s.n + 1;  
    if (! empty (s.L)) {  
        p = removeProc (s.L);  
        unblock (p);  
    }  
    <exit>  
}
```

- Give up mutual exclusion before
- Reestablish after?

# Race Condition!

## Implementation 2

```
wait (sem s) {  
    <entry>  
    if s.n ≤ 0 {  
        addProc (me, s.L);  
        <exit>  
        block (me);  
        <entry>  
    }  
    s.n = s.n - 1;  
    <exit>  
}
```



```
signal (sem s) {  
    <entry>  
    s.n = s.n + 1;  
    if (! empty (s.L)) {  
        p = removeProc (s.L);  
        unblock (p);  
    }  
    <exit>  
}
```

- What if context switch happens just after unblocking but before <entry>